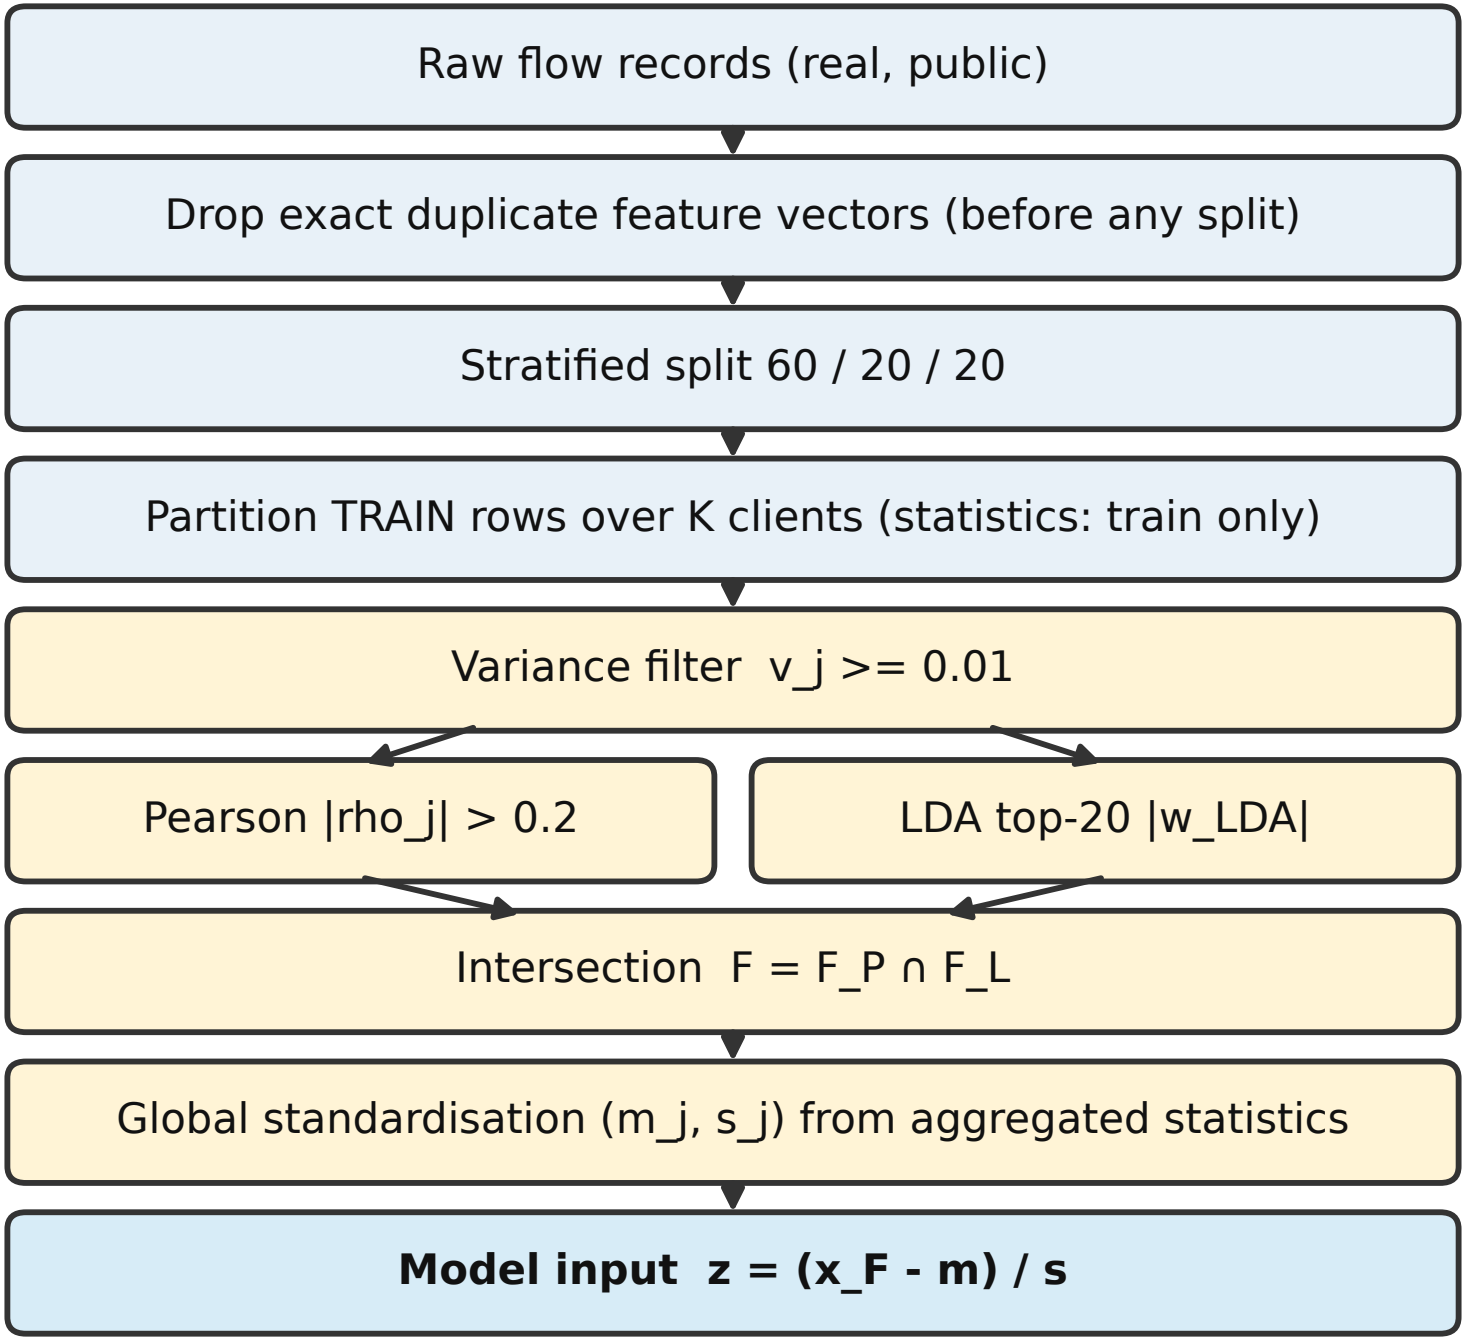
Raw flow records (real, public)
Drop exact duplicate feature vectors (before any split)
Stratified split 60 / 20 / 20
Partition TRAIN rows over K clients (statistics: train only)
Variance filter v_j >= 0.01
Pearson |rho_j| > 0.2
LDA top-20 |w_LDA|
Intersection F = F_P ∩ F_L
Global standardisation (m_j, s_j) from aggregated statistics
Model input z = (x_F - m) / s