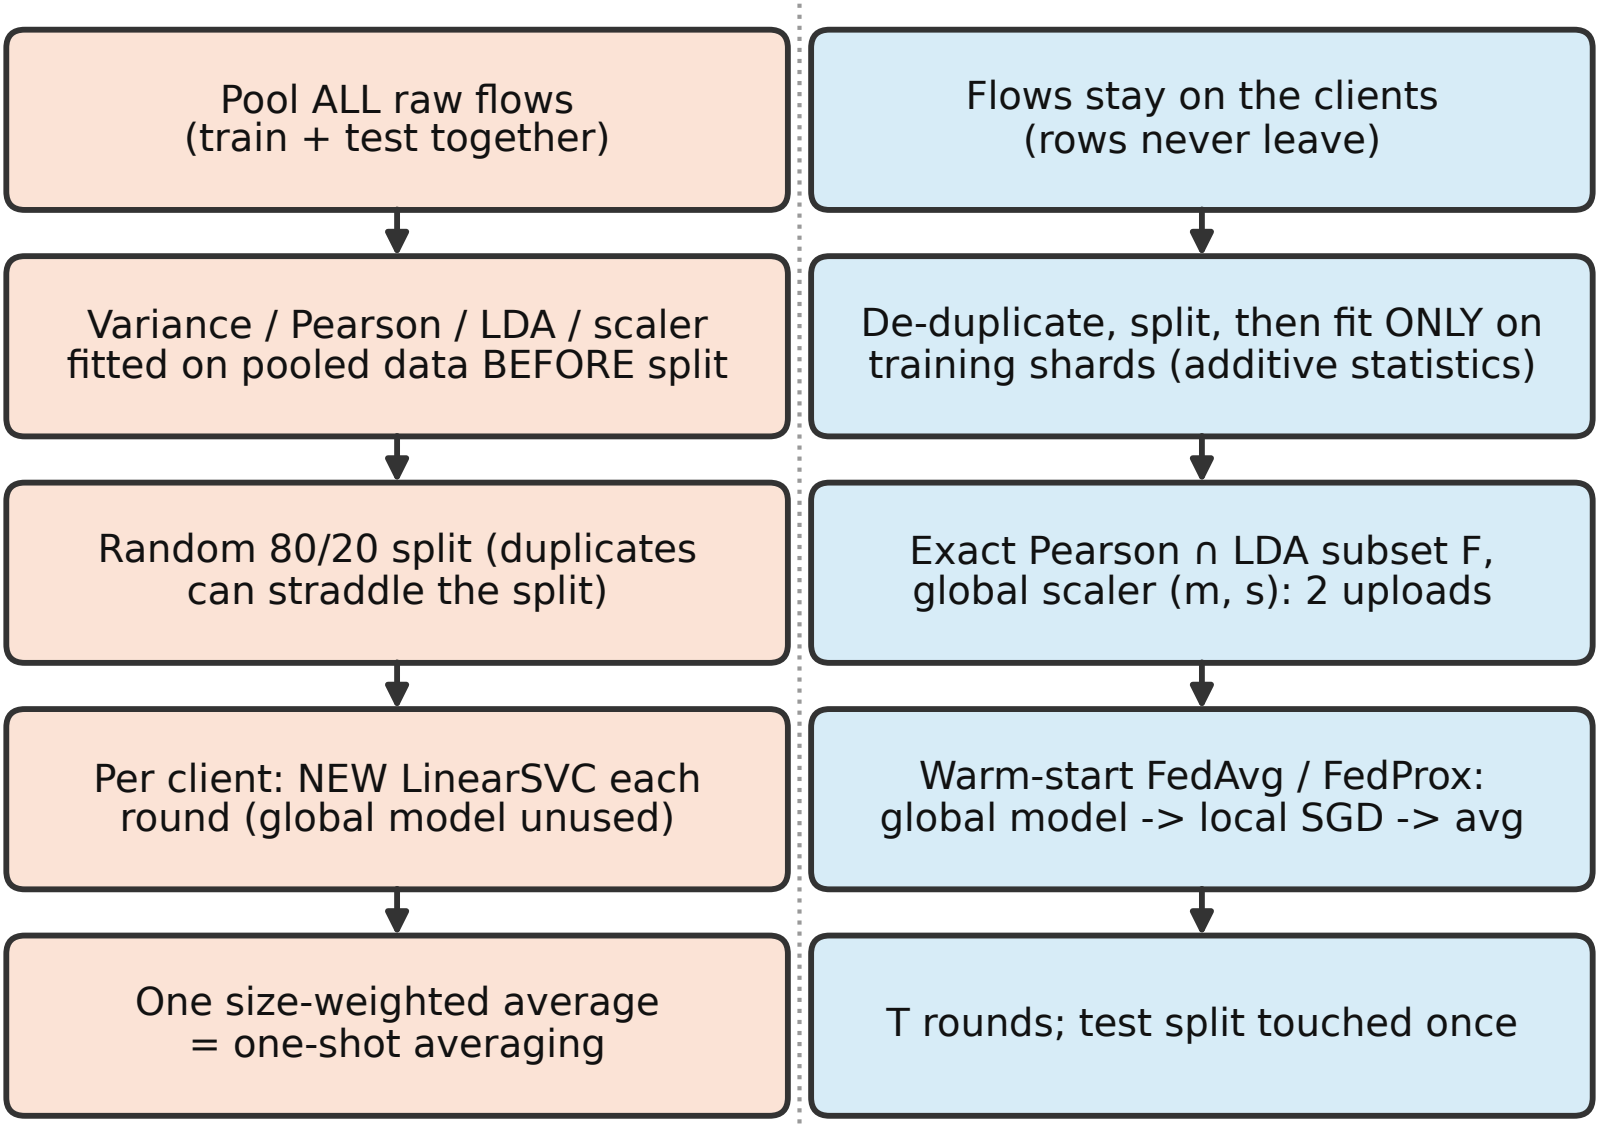
Senior pipeline (as released)
Pool ALL raw flows
(train + test together)
Variance / Pearson / LDA / scaler
fitted on pooled data BEFORE split
Random 80/20 split (duplicates
can straddle the split)
Per client: NEW LinearSVC each
round (global model unused)
One size-weighted average
= one-shot averaging
Proposed FedStat-FS
Flows stay on the clients
(rows never leave)
De-duplicate, split, then fit ONLY on
training shards (additive statistics)
Exact Pearson ∩ LDA subset F,
global scaler (m, s): 2 uploads
Warm-start FedAvg / FedProx:
global model -> local SGD -> avg
T rounds; test split touched once